Telecom Churn Case Study

Bhagabati Mishra Akash Chaterjee Biswanath Nanda

Description	Acronyms	
Customer phone numl	MOBILE_NUMBER	0
Telecom circle area to which the customer below	CIRCLE_ID	1
Local calls - within same telecom cir	LOC	2
STD calls - outside the calling cir	STD	3
Incoming ca	IC	4
Outgoing ca	OG	5
Operator T to T, i.e. within same operator (mo	T2T	6
Operator T to other operator mot	T2M	7
Operator T to other operator fixed I	T2O	8
Operator T to fixed lines of	T2F	9
Operator T to it's own call cen	T2C	10
Average revenue per u	ARPU	11
Minutes of usage - voice ca	MOU	12
Age on network - number of days the customer	AON	13
All kind of calls within the same operator ne	ONNET	14
All kind of calls outside the operator T netwo	OFFNET	15
Indicates that customer is in roaming zone du	ROAM	16
Special ca	SPL	17
ISD ca	ISD	18
Recha	RECH	19

Case Introduction & Data

- Customers usually do not decide to switch to another competitor instantly, but rather over a period (this is especially applicable to high-value customers). In churn prediction, we assume that there are **three phases** of the customer lifecycle-
 - The 'good' phase: In this phase, the customer is happy with the service and behaves as usual.
 - The 'action' phase: The customer experience starts to sore in this
 phase, for e.g. he/she gets a compelling offer from a competitor,
 faces unjust charges, becomes unhappy with service quality etc. In
 this phase, the customer usually shows different behaviour than in
 the 'good' months. Also, it is crucial to identify high-churn-risk
 customers in this phase, since some corrective actions can be
 taken at this point (such as matching the competitor's
 offer/improving the service quality etc.)
 - The 'churn' phase: In this phase, the customer is said to have churned.
 - Provided data has 999,999 rows and 226 columns
 - 179- float64
 - 35- int64
 - 12- object

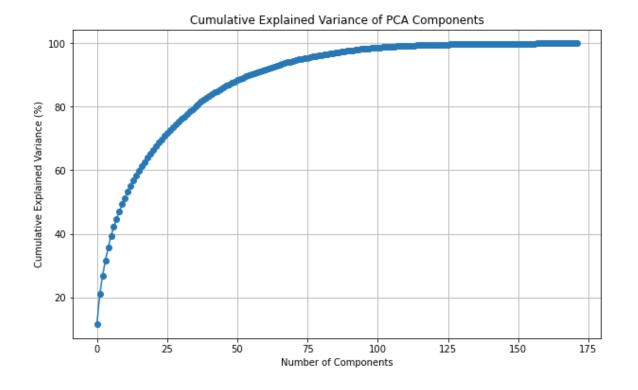
Finding Valuable Customers

- As prescribed by the user, onboarding new customers are expensive than retaining. Hence, predicting potential churning customers are of upmost importance.
- However, it makes more sense to focus on valuable customers. These are the customers who generate revenue for the company frequently.

```
# Calculate the total data-recharge and total recharge amount amount for June and July
   churn['total_data_rech_6'] = churn.total_rech_data_6 * churn.av_rech_amt_data_6
   churn['total_data_rech_7'] = churn.total_rech_data_7 * churn.av_rech_amt_data_7
   churn['amt_data_6'] = churn.total_rech_amt_6 + churn.total_data_rech_6
   churn['amt_data_7'] = churn.total_rech_amt_7 + churn.total_data_rech_7
   # calculate average recharge done by customer in June and July
10 churn['av_amt_data_6_7'] = (churn.amt_data_6 + churn.amt_data_7)/2
 1 # look at the 70th percentile recharge amount
   print("Recharge amount at 70th percentile: {0}".format(churn.av_amt_data_6_7.quantile(0.7)))
   # Keeping customers who have recharged their mobiles with more than or equal to 70th percentile amount
   churn_filtered = churn.loc[churn.av_amt_data_6_7 >= churn.av_amt_data_6_7.quantile(0.7), :]
   churn_filtered = churn_filtered.reset_index(drop=True)
 9 churn filtered.shape
Recharge amount at 70th percentile: 478.0
(30001, 201)
 1 # Removing variables created to filter high-value customers
   churn_filtered = churn_filtered.drop(['total_data_rech_6', 'total_data_rech_7',
                                          'amt_data_6', 'amt_data_7', 'av_amt_data_6_7'], axis=1)
   churn_filtered.shape
```

Steps taken while building the model

- 1. Standard Scaler: This step scales the features of the dataset to have a mean of 0 and a standard deviation of 1.
- 2. PCA (Principal Component Analysis): This step performs dimensionality reduction by projecting the data onto a lower-dimensional subspace while retaining most of the variance in the data.



Evaluating the model

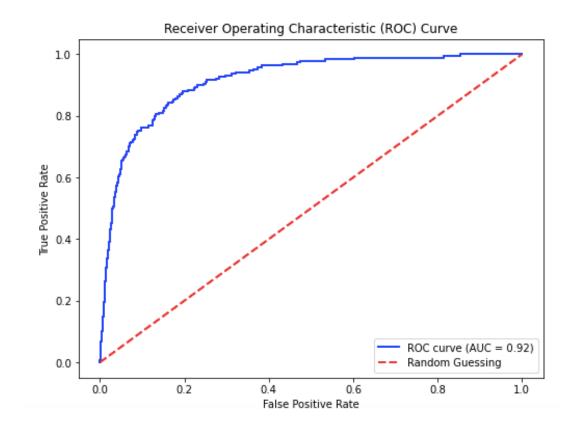
Confusion Matrix:

- True Positives (TP): These are the correctly predicted churn instances. These are the customers
 who were predicted to churn and actually churned.
- True Negatives (TN): These are the correctly predicted non-churn instances. These are the customers who were predicted not to churn and actually did not churn.
- False Positives (FP): These are the incorrectly predicted churn instances. These are the customers
 who were predicted to churn but actually did not churn. This is also known as Type I error or false
 alarms.
- False Negatives (FN): These are the incorrectly predicted non-churn instances. These are the
 customers who were predicted not to churn but actually churned. This is also known as Type II
 error or missed opportunities.

Interpretation:

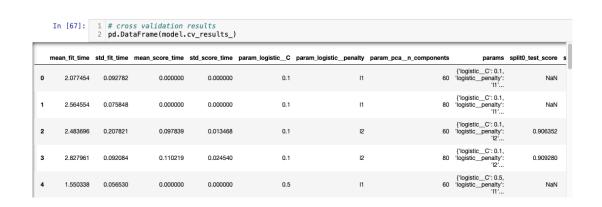
- Accuracy: It measures how often the classifier makes the correct prediction, and it's calculated as (TP + TN) / (TP + TN + FP + FN).
- Precision: It measures the accuracy of positive predictions, and it's calculated as TP / (TP + FP). It
 tells the proportion of correctly predicted churn instances among all instances predicted as churn.
- Recall (Sensitivity): It measures the proportion of actual churn instances that were correctly predicted by the classifier, and it's calculated as TP / (TP + FN).
- Specificity: It measures the proportion of actual non-churn instances that were correctly predicted by the classifier, and it's calculated as TN / (TN + FP).

In the context of telecom churn prediction, minimizing false negatives (missed churn opportunities) is typically more critical than minimizing false positives (incorrectly identifying non-churners as churners). Therefore, we want to focus on improving recall while keeping a reasonable level of precision.



Peforming Cross-Validation

- 1. Logistic Regression Initialization: Logistic regression is initialized. The class_weight parameter is set to handle class imbalance, giving higher weight to the minority class (class 1, churn) compared to the majority class (class 0, non-churn).
- **2. Pipeline Creation**: A pipeline is created with the following steps:
 - 1. StandardScaler: Scales the features of the dataset.
 - 2. PCA: Performs dimensionality reduction.
 - 3. Logistic Regression: Classifies the data.
- **3. Hyperparameter Space**: Define a dictionary params specifying the hyperparameter space for grid search. It includes options for the number of components in PCA (n_components), the regularization parameter C, and penalty penalty for logistic regression.
- 4. Cross-Validation Folds: Create 5 folds for cross-validation using StratifiedKFold. It ensures that each fold maintains the same class distribution as the original dataset.
- **5. Grid Search**: GridSearchCV is initialized with the pipeline, cross-validation folds, hyperparameter space, scoring metric (ROC AUC), and settings for parallel execution (n_jobs=-1 for utilizing all available CPU cores). This object is ready for fitting to find the best hyperparameters for the model.



Using the final model that has is more stable and decent sensitivity score, the firm can focus on the customers classified as potential churners and take proactive actions.

Final Model

```
# predict churn on test data
y_pred = model.predict(X_test)

# create onfusion matrix
cm = confusion_matrix(y_test, y_pred)
print(cm)

# check sensitivity and specificity
sensitivity, specificity, _ = sensitivity_specificity_support(y_test, y_pred, average='binary')
print("Sensitivity: \t", round(sensitivity, 2), "\n", "Specificity: \t", round(specificity, 2), sep='')

# check area under curve
y_pred_prob = model.predict_proba(X_test)[:, 1]
print("ROC: \t", round(roc_auc_score(y_test, y_pred_prob),2))
[[6394 324]
[ 72 1181]
```

[72 118]]
Sensitivity: 0.62
Specificity: 0.95
ROC: 0.9